

# Importance Sampling

May 16, 2021

```
[1]: import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns; sns.set()
import scipy
```

## 1 Importance Sampling

**Problem** Given a random variable  $X$  following a distribution with  $p$  as its probability density function (pdf). We want to estimate the expectation  $\mathbb{E}_p(f(X))$  where  $f$  is an arbitrary function.

### Method (Importance Sampling)

One way to estimate is to use the Monte Carlo approximation method, in which we sample values  $x_1, x_2, \dots, x_N$  from  $p$  and make the estimation

$$\mathbb{E}_p(f(X)) = \int_{-\infty}^{\infty} f(x)p(x)dx \approx \frac{1}{N} \sum_{i=1}^N f(x_i).$$

However,  $p$  may be hard to sample from. Instead of doing so, we sample from a proposal density  $q$  and estimate based on the values.

The key idea of Importance Sampling is that

$$\int_{-\infty}^{\infty} f(x)p(x)dx = \int_{-\infty}^{\infty} f(x)\frac{p(x)}{q(x)}q(x)dx = \mathbb{E}_q \left[ f(X)\frac{p(X)}{q(X)} \right]$$

which can be approximated in the same way as above: sampling values  $x_1, x_2, \dots, x_N$  from  $q$  and computing

$$\mathbb{E}_q \left[ f(X)\frac{p(X)}{q(X)} \right] = \frac{1}{N} \sum_{i=1}^N \left[ f(x_i)\frac{p(x_i)}{q(x_i)} \right].$$

The name “importance” stems from the ratio  $p(x_i)/q(x_i)$  being considered the importance weights of the sample  $x_i$  (which makes sense).

```
[2]: def importance_sampling(f, p_pdf, q_pdf, q_sample, N, return_history=False):
    x = q_sample(N)
    f = f(x) * p_pdf(x) / q_pdf(x)
    history = np.cumsum(f) / np.arange(1, N+1, 1)
```

```

if return_history:
    return history[-1], history
else:
    return history[-1]

```

**Example (Expectation of the Beta Distribution)** A random variable  $X$  follows the Beta distribution  $\text{Beta}(a, b)$ , calculate its expectation.

The PDF of  $X$  is:

$$p(x; a, b) = \frac{(a + b - 1)!}{(a - 1)!(b - 1)!} x^{a-1} (1 - x)^{b-1}$$

when  $x \in [0, 1]$  and 0 otherwise.

To calculate the expectation is to calculate  $\mathbb{E}(f(X))$  with  $f(x) = x$ .

```

[3]: def beta_pdf(x, a, b):
    return scipy.stats.beta.pdf(x, a, b)


def beta_expectation(a, b):
    return a / (a + b)

[4]: As = [2, 6, 3]
Bs = [8, 6, 1]

fig = plt.figure(figsize=(8, 4), dpi=120)
x = np.linspace(start=0, stop=1, num=1000)
for A, B in zip(As, Bs):
    f = beta_pdf(x, A, B)
    p = plt.plot(x, f, label=f'$a = {A}, b = {B}$')

    e = beta_expectation(A, B)
    plt.plot(e, 0, 'o', color=p[0].get_color())
plt.title('PDF of $X \sim \text{Beta}(a, b)$')
plt.legend()
plt.show()

```



We first use the uniform distribution on  $[0, 1]$  as the proposal density. The PDF is  $q(x; a, b) = 1.0, \forall x \in [a, b]$  and 0 otherwise.

```
[5]: def uniform_pdf(x, a=0.0, b=1.0):
      return 1.0 * (a <= x) * (x <= b)

      def uniform_sample(N, a=0.0, b=1.0):
          return np.random.rand(N) * (b - a) + a
```

```
[6]: a1 = 2.0
      b1 = 8.0
      N1 = 5000000

      f = lambda x: x
      p_pdf = lambda x: beta_pdf(x, a1, b1)
      q_pdf = uniform_pdf
      q_sample = uniform_sample

      u1, h1 = importance_sampling(f, p_pdf, q_pdf, q_sample, N=N1,
      ↪return_history=True)
```

```
[7]: r1 = beta_expectation(a1, b1)

      print('Estimated value:', u1)
      print('Real value', r1)

      fig = plt.figure(figsize=(6, 3), dpi=120)
```

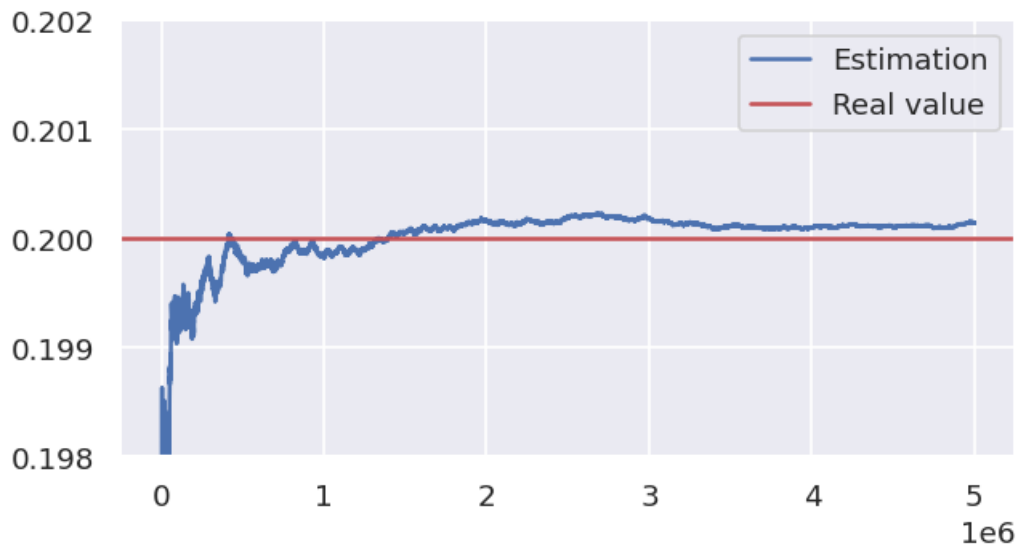
```

step = 100
x = range(0, len(h1), step)
sns.lineplot(x=x, y=h1[x], label='Estimation')
plt.axhline(r1, color='r', label='Real value')
plt.ylim(r1 - 0.002, r1 + 0.002)
plt.legend()
plt.show()

```

Estimated value: 0.2001462734509228

Real value 0.2



```

[8]: a2 = 3.0
     b2 = 1.0
     N2 = 5000000

     f = lambda x: x
     p_pdf = lambda x: beta_pdf(x, a2, b2)
     q_pdf = uniform_pdf
     q_sample = uniform_sample

     u2, h2 = importance_sampling(f, p_pdf, q_pdf, q_sample, N=N2,
                                ↪return_history=True)

```

```

[9]: r2 = beta_expectation(a2, b2)

     print('Estimated value:', u2)
     print('Real value', r2)

```

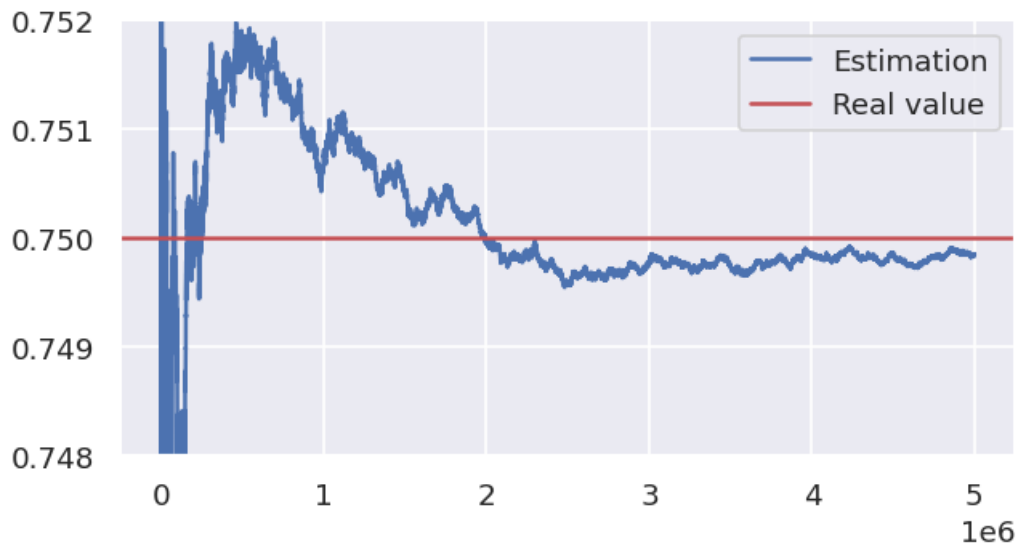
```

fig = plt.figure(figsize=(6, 3), dpi=120)
step = 100
x = range(0, len(h2), step)
sns.lineplot(x=x, y=h2[x], label='Estimation')
plt.axhline(r2, color='r', label='Real value')
plt.ylim(r2 - 0.002, r2 + 0.002)
plt.legend()
plt.show()

```

Estimated value: 0.749845498264427

Real value 0.75



We then use the normal distribution. The PDF is

$$q(x; \mu, \sigma) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(x - \mu)^2}{2\sigma^2}\right), \forall x \in \mathbb{R}.$$

```

[10]: def normal_pdf(x, mu, sigma):
        return scipy.stats.norm.pdf(x, mu, sigma)

def normal_sample(N, mu, sigma):
    return np.random.randn(N) * sigma + mu

```

```

[11]: a1 = 2.0
b1 = 8.0
N1 = 5000000

f = lambda x: x
p_pdf = lambda x: beta_pdf(x, a1, b1)

```

```

mu1, sigma1 = 0.2, 0.2
q1_pdf = lambda x: normal_pdf(x, mu1, sigma1)
q1_sample = lambda N: normal_sample(N, mu1, sigma1)

mu2, sigma2 = 0.8, 0.2
q2_pdf = lambda x: normal_pdf(x, mu2, sigma2)
q2_sample = lambda N: normal_sample(N, mu2, sigma2)

u3, h3 = importance_sampling(f, p_pdf, q1_pdf, q1_sample, N=N1,
    ↪return_history=True)
u4, h4 = importance_sampling(f, p_pdf, q2_pdf, q2_sample, N=N1,
    ↪return_history=True)

```

```

[12]: r1 = beta_expectation(a1, b1)

print('Estimated value 1:', u3)
print('Estimated value 2:', u4)
print('Real value', r1)


fig = plt.figure(figsize=(6, 3), dpi=120)
step = 100
x = range(0, len(h3), step)
sns.lineplot(x=x, y=h3[x], label='Estimation 1')
sns.lineplot(x=x, y=h4[x], label='Estimation 2')
plt.axhline(r1, color='r', label='Real value')
plt.ylim(r1 - 0.002, r1 + 0.002)
plt.legend()
plt.show()

```

```

Estimated value 1: 0.2000085412182431
Estimated value 2: 0.1984876196104516
Real value 0.2

```




The proposal  $\mathcal{N}(0.2, 0.2^2)$  seems to perform better than  $\mathcal{N}(0.8, 0.2^2)$ . What seems to be the problem here?

```
[13]: fig = plt.figure(figsize=(8, 4), dpi=120)

x = np.linspace(start=0, stop=1, num=1000)
f = x
p = beta_pdf(x, 2, 8)
q1 = normal_pdf(x, 0.2, 0.2)
q2 = normal_pdf(x, 0.8, 0.2)
plt.plot(x, p, label='True PDF p', linestyle='dashed')
plt.plot(x, q1, label='Reference q1')
plt.plot(x, q2, label='Reference q2')
plt.plot(x, np.abs(f) * p, label='|f|*p')

plt.legend()
plt.show()
```



Intuitively, when we estimate (the expectation of)  $f(X)$  we should be paying attention to the areas where  $x$  has a higher chance of contributing **and**  $f(x)$  has high (absolute) values so that it contributes more to the expectation. It is suggested we pick a proposal  $q$  so that the probability is dense ( $q(x)$  is high) where  $|f(x)|p(x)$  is large. From the graph, the first normal distribution is a better choice and the effect can be seen from the example above.

Note that it is not true that  $q$  should closely resemble  $p$  or has similar peaks.

## 2 Self-normalizing weights

Importance Sampling is possible even for un-normalized target distribution, in other words,  $\tilde{p}$  so that  $p(x) = \tilde{p}(x)/z_p$ . It is even possible when we have un-normalized proposal distribution, in other words,  $\tilde{q}$  so that  $q(x) = \tilde{q}(x)/z_q$ .

$z_p, z_q$  are the normalizing constants of  $p$  and  $q$  respectively. This normalizing constant is usually to ensure the PDF integrates to 1, meaning that  $z_p = \int_{-\infty}^{\infty} \tilde{p}(x)dx$  and same for  $z_q$ .

Same as before, we have

$$\int_{-\infty}^{\infty} f(x)p(x)dx = \int_{-\infty}^{\infty} f(x)\frac{p(x)}{q(x)}q(x)dx = \int_{-\infty}^{\infty} f(x)\frac{\tilde{p}(x)}{z_p}\frac{z_q}{\tilde{q}(x)}q(x)dx = \frac{z_q}{z_p}\mathbb{E}_q\left[f(X)\frac{\tilde{p}(X)}{\tilde{q}(X)}\right]$$

The difference is the ratio  $z_q/z_p$ . We can approximate this ratio in parallel by noticing that

$$\frac{z_p}{z_q} = \frac{1}{z_q} \int_{-\infty}^{\infty} \tilde{p}(x)dx = \frac{1}{z_q} \int_{-\infty}^{\infty} \frac{\tilde{p}(x)}{\tilde{q}(x)}\tilde{q}(x)dx = \int_{-\infty}^{\infty} \frac{\tilde{p}(x)}{\tilde{q}(x)}\frac{\tilde{q}(x)}{z_q}dx = \int_{-\infty}^{\infty} \frac{\tilde{p}(x)}{\tilde{q}(x)}q(x)dx = \mathbb{E}_q\left[\frac{\tilde{p}(X)}{\tilde{q}(X)}\right]$$

Once again using Monte-Carlo approximation, we sample values  $x_1, x_2, \dots, x_N$  from  $q$  and compute

$$\frac{z_p}{z_q} \approx \sum_{i=1}^N \frac{\tilde{p}(x_i)}{\tilde{q}(x_i)}$$



In summary, to estimate  $\mathbb{E}_p(f(X))$  for  $x \sim p(x) = \tilde{p}(x)/z_p$  based on a proposed distribution  $q(x) = \tilde{q}(x)/z_q$  (which we can easily sample from), we sample  $x_1, x_2, \dots, x_N$  and compute

$$\mathbb{E}_p(f(X)) \approx \frac{\sum_{i=1}^N f(x_i) \frac{\tilde{p}(x_i)}{\tilde{q}(x_i)}}{\sum_{i=1}^N \frac{\tilde{p}(x_i)}{\tilde{q}(x_i)}}$$

```
[14]: def importance_sampling_unnorm(f, p_pdf, q_pdf, q_sample, N,
    ↪return_history=False):
    x = q_sample(N)
    w = p_pdf(x) / q_pdf(x)
    fw = f(x) * w
    history = np.cumsum(fw) / np.cumsum(w)

    if return_history:
        return history[-1], history
    else:
        return history[-1]
```

### Example

The un-normalized PDF of the Beta distribution is

$$\tilde{p}(x; a, b) = x^{a-1}(1-x)^{b-1}$$

when  $x \in [0, 1]$  and 0 otherwise.

The un-normalized PDF of the Normal distribution is

$$\tilde{q}(x; \mu, \sigma) = \exp\left(-\frac{(x-\mu)^2}{2\sigma^2}\right), \forall x \in \mathbb{R}.$$

```
[15]: def beta_unnorm_pdf(x, a, b):
    p = x ** (a - 1) * (1 - x) ** (b - 1)
    p[x > 1] = 0.0
    p[x < 0] = 0.0
    return p

def normal_unnorm_pdf(x, mu, sigma):
    return np.exp(-(x - mu)**2 / (2 * sigma**2))
```

```
[16]: a1 = 2.0
    b1 = 8.0
    N1 = 5000000

    f = lambda x: x
    p_pdf = lambda x: beta_unnorm_pdf(x, a1, b1)

    mu1, sigma1 = 0.2, 0.2
```

```

q1_pdf = lambda x: normal_unnorm_pdf(x, mu1, sigma1)
q1_sample = lambda N: normal_sample(N, mu1, sigma1)

mu2, sigma2 = 0.8, 0.2
q2_pdf = lambda x: normal_unnorm_pdf(x, mu2, sigma2)
q2_sample = lambda N: normal_sample(N, mu2, sigma2)

u3, h3 = importance_sampling_unnorm(f, p_pdf, q1_pdf, q1_sample, N=N1,
    ↪return_history=True)
u4, h4 = importance_sampling_unnorm(f, p_pdf, q2_pdf, q2_sample, N=N1,
    ↪return_history=True)

```

```

[17]: r1 = beta_expectation(a1, b1)

print('Estimated value 1:', u3)
print('Estimated value 2:', u4)
print('Real value', r1)

fig = plt.figure(figsize=(6, 3), dpi=120)
step = 100
x = range(0, len(h3), step)
sns.lineplot(x=x, y=h3[x], label='Estimation 1')
sns.lineplot(x=x, y=h4[x], label='Estimation 2')
plt.axhline(r1, color='r', label='Real value')
plt.ylim(min(h3[-1], h4[-1]) - 0.005, max(h3[-1], h4[-1]) + 0.005)
plt.legend()
plt.show()

```

Estimated value 1: 0.1999546313998175  
 Estimated value 2: 0.19971797642382474  
 Real value 0.2

